

Project - Parallel Wave Function Collapse Terrain Generation using CUDA and MPI

COMP 4901Q: High Performance Computing (HPC) Spring 2022

CHEUK NAM TSANG

HKUST

cntsangaa@connect.ust.hk

1. Introduction

Procedural content generation means methods of generating content based on set rules. Commonly used in video games such as Minecraft and Crypt of the NecroDancer, procedural generation of game environment can allow game developers to design a set of rules that can generate unique permutations of environment for each player without manually designing all of the environment features.

Wave Function Collapse Algorithm is one of the tools for generative design of video game level environment such as terrains [Moreau 2020]. The algorithm is similar to a game of Sudoku. In an empty 9x9 board, all tiles can be all of the 9 numbers. When one number is placed, tiles on the same row, same column or same 3x3 square can no longer be assigned the same number.

Following are the procedures for the algorithm:

1. Initialize the input arena of size $N \times N$ individual cells
2. Create an array w to list all valid states
3. Initialize a boolean array for each cells with has the same size as array w the that indicates whether the state is valid or forbidden.
4. Pick an arbitrary cell and randomly set it to a valid state. Usually, cells with low number of possible states are prioritized to prevent conflicts.
5. Recalculate valid states for all cells.
6. Repeat from step 4 until all cells has defined state.

2. Related works

In 2019, a team of students attempted to parallelize the algorithm using different approaches [Orlowski and Lee 2019]. They used OpenMP and P-Threads in combination with Queues. Their results showed that the algorithm can be parallelized with relatively high scaling efficiency. However, it also revealed that an observer-based approach with sequential queue for execution shows the most speedup when compared to parallelizing the work on all cells.

In this project, I aim to investigate parallelizing the wave function collapse algorithm using CUDA since CUDA has many more cores compared to multiprocessing on CPU, the brute-force approach on CUDA may still be beneficial over CPU with a faster algorithm.

Since the next iteration is based on changes from adjacent cells, this is very similar to Conway's game of life [Games 1970], where the livelihood of cells in the next iteration are entirely based on the cells surrounding it in the current iteration. Implementation of this game using image convolution and

utilizing GPU shader processing units has been proven possible [Ru 2017]. This project also tries to utilize multiple GPU by using MPI, which can further speed up the process.

3. Target Platform

This project targets a multi-GPU cluster, and the problem size will need to be big in order to benefit from parallelization. The cluster is expected to be able to communicate with each other using MPI and each servers should have CUDA capable GPU so that the CUDA program can be executed.

4. Application Objectives

This program will not have a graphical user interface or graphical representation of the results other than being able to export to bitmap.

For simplicity, all cells can be allocated an integer from 0 to 9 and only one rule will apply. Adjacent cells can have a difference of at most 1, meaning cells surrounding a cell with state 8 can only has the value 7, 8 and 9. This also makes propagating the states faster since Manhattan distance can be used to limit the range of propagation.

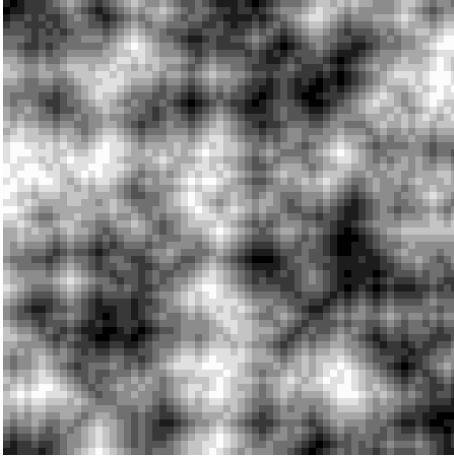


Figure 1. Generated height map.

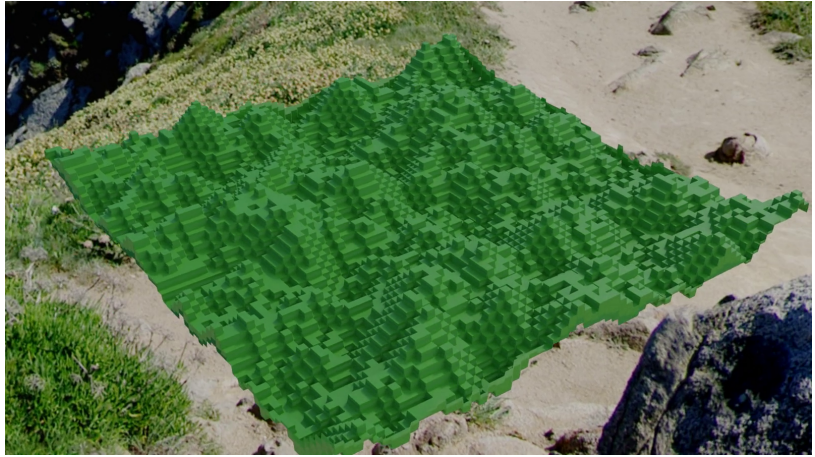


Figure 2. Generated terrain rendered in a 3D environment.

This rule set simulates generation of mountain terrain in video games where no steep cliffs are present and the map has gradually changing attitudes. As can be seen in the images above, this algorithm aims to programmatically generate organic-seeming terrains.

5. Baseline Implementation

The serial program will first initialize the map with user specified size, and then start the wave function collapse algorithm. After collapsing state of each individual cell, All neighbors with more than one valid states will need to be updated with the following formula:

$$lowerBound = \max(lowerBound, collapsed.value - ManhattanDistance(collapsed, this))$$

$$upperBound = \min(upperBound, collapsed.value + ManhattanDistance(collapsed, this))$$

Since there are only 10 states, only cells with a Manhattan distance of 9 or below will need to have their states updated. This propagation optimization reduces the complexity but still consumes majority

of the runtime in previous studies. A cell is randomly chosen to has its state collapsed until all cells has a definite state.

Once the execution has completed, the solution will be verified against the rule set. In addition, user can specify that the generated result to be saved to a bitmap image file so that it can be read by other software for visualization, as can be seen in figure 1 and 2.

6. Parallel Implementation

6.1 Theoretical Background

The specific set of rules implemented in this project which limits the side effects to neighboring cells. This rule set is common in game terrain generation but it is also popular because of its locality benefits. By locality, it means that non-determinate cells encapsulated by determinate cells will not be restricted further by cells outside of the walls. It also means that a map can be divided into smaller sections and allow each processes to handle their local area without depending their calculation on results from other processes. This project will try to take advantage of this task division method to spread the task to all processes in the MPI-connected system.

Moreover, the locality also limits the number of cells to update since their restrictions will be based on cells with Manhattan distance of 9 or below, which means that all cells with the newly defined cell can be updated simultaneously without depending on calculations from the previous updates. This project will try to take advantage of this and process all cell updates using parallel CUDA cores.

6.2 Implementation Details

First, the map is divided into strips of height 9. The top horizontal line of each strip are independent from the top horizontal line of other strips because the Manhattan distance is above the limit of the propagation in this rule set. Thus, the first row of each strip are independently initialized on different processes in the MPI system. Moreover, collapsing cells in a straight line can delay propagation of non-forbidden states in neighboring cells to after the entire edge is initialized.

Then, each process will propagate their top edge to the process responsible for the strip above. This allows for each process to have a localized copy of the problem with a smaller problem size, yet all of them are compatible since the connecting edges are synchronized. Each process have to propagate the constraints from both top and bottom edges before starting to fill in the remaining cells accordingly.

MPI processes performs calculation on CUDA-enabled GPU for parallel constraints propagation. This can be done because constraint propagation are fitting for SIMD architecture. This is akin to a kernel in convolution but has a diamond shape because it is created by limiting Manhattan distance. When rotated 45 degree, the kernel can fit inside a grid of n by $2n+1$ where n is the maximum Manhattan distance. However, a larger grid of threads can be initialized instead so as to simplify the indexing calculation.

Lastly, the master rank in the MPI system will collect the strips and combine them into one image which can then be verified or optionally saved to a bitmap file.

The feature to export the resultant map into a 2D image allowed visual confirmation of the correctness of the program on top of the verification implemented that only checks whether the map complies with the rules. This is an important feature not only to allow the computation result to be saved, but also debug versions of the program where a matrix of all zero was returned which can pass the verification but is not the desired result.

7. Performance Scaling

With both MPI and CUDA independently implemented, the project can test and verify the performance scaling on various configurations. The three testable configurations are CUDA, MPI and MPI+CUDA.

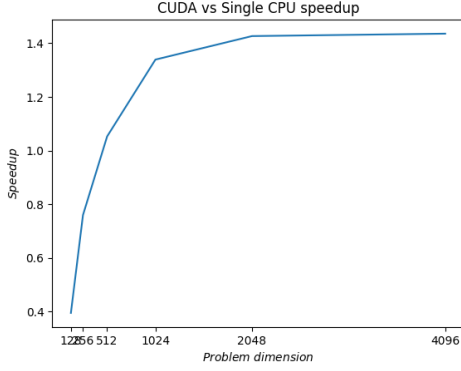


Figure 3. Speedup of using CUDA on different problem sizes.

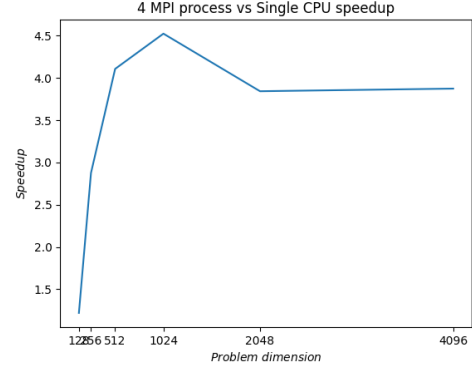


Figure 4. Speedup of using 4 MPI processes on different problem sizes.

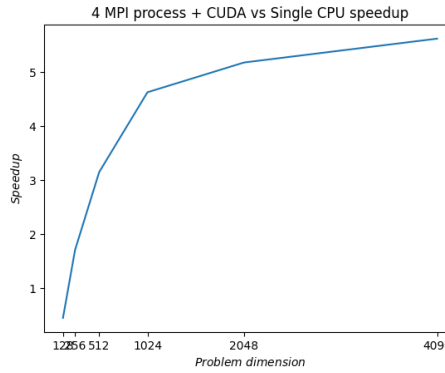


Figure 5. Speedup of using 4 MPI processes with CUDA on different problem sizes.

As can be seen from the graphs above, the speedup of CUDA is disappointing topping at around 42% speedup on from the sequential version. This is a result of the overhead in initializing the GPU's global memory and the lack of locality during the propagation. Moreover, each iteration of collapse and propagation are dependent on results from the last iteration, limiting the number of warps active on the GPU to one. This is an oversight in the initial vision of this project since propagation was thought to be similar to convolution which CUDA is known to have an advantage over the CPU. The speedup managed in this implementation is mainly due to the threading and allowing all propagation to be done without looping.

On the MPI performance, it is apparent that the overhead of collecting the final map to the CPU is significant in smaller problem sizes. Curiously, at problem size of around 1024, the performance gain is greater than the number of processes allocated for the task. This could be the result of splitting the map into sections that limited divergence and increased the rate of wave function collapse which can reduce the number of calls to the random number generator.

When combined, the speedup diminishes at 5.5 times for 4 CUDA-aware MPI processes, and shows an optimal problem size of at least 1024.

On performance increase as the number of processors increases, the CPU-only implementation is compared against single CPU implementation, and the MPI+CUDA implementation is compared against single GPU implementation. Both are tested with a constant problem size of 2048.

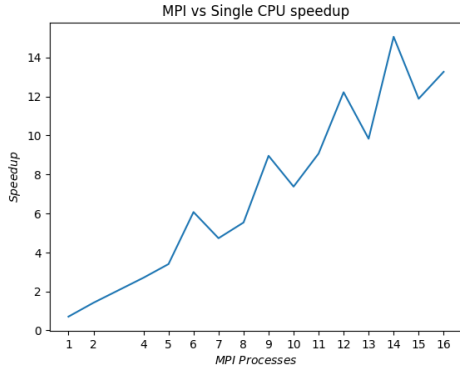


Figure 6. Speedup of MPI on different number of processes.

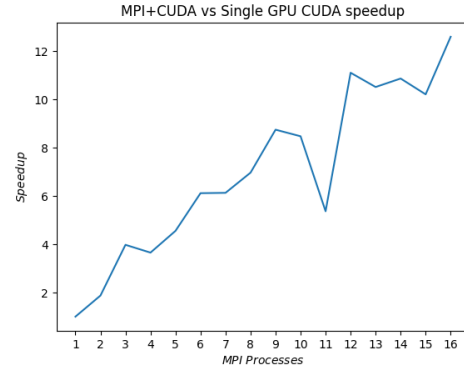


Figure 7. Speedup of MPI+CUDA on different number of processes.

As shown above, the scaling efficiency of this MPI-oriented algorithm is high with almost linear scaling. And the CUDA-aware version shows that this the MPI communication plan can work in conjunction with other parallelization techniques. However, there exist a downward spike when allocated 11 CUDA-aware MPI processes. This could be a result of uneven problem sizes per process that causes more redundant calculation in boundary cases.

In general, using power of twos as number of processes are favorable in terms of scaling efficiency and reducing overhead.

8. Limitations

This implementation of the Wave Function Collapse algorithm has two main limitations. First is the restrictive rules. Since the algorithm is designed to optimize for the specific set of rules, it cannot be easily applied to speed up other rule sets. In other words, the proposed solutions are not general to all Wave Function Collapse algorithms.

The second weakness of this algorithm is the oversight in how CUDA is not entirely suitable for this problem. In more general implementation of the Wave Function Collapse algorithm, processes are highly divergent and very dependent on results from previous iterations. The simplified set of rules allowed CUDA GPUs to accelerate part of the propagation but the lack of concurrent warps limits the effectiveness of GPU utilization.

9. Algorithm Tuning

Initial version of the program attempted to collapse the wave of cells in sequential order, each row will be collapsed from left to right, then the next row. Although the algorithm can return a correct solution, this poses the problem of the resultant image showing diagonal patterns since the constraints of states consistently start from one direction to another.

To enable pseudo random ordering without repeatedly check on cells or keep track of all cells in a queue, modulo addition of prime number is used. As suggested on similar discussion on Stack Overflow [Lindley 2013], modulo addition of a prime number larger than dimension of the array can jump the indices without repeating or missing any in a full set of iterations.

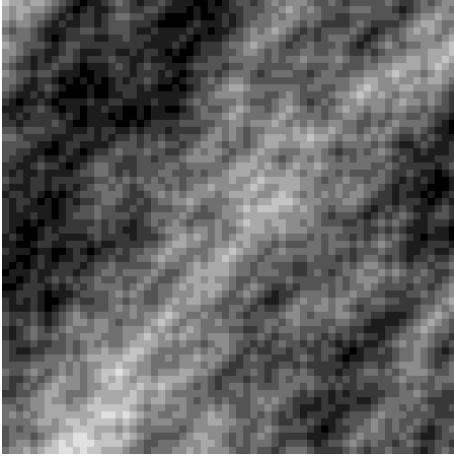


Figure 8. Collapsed in sequential order

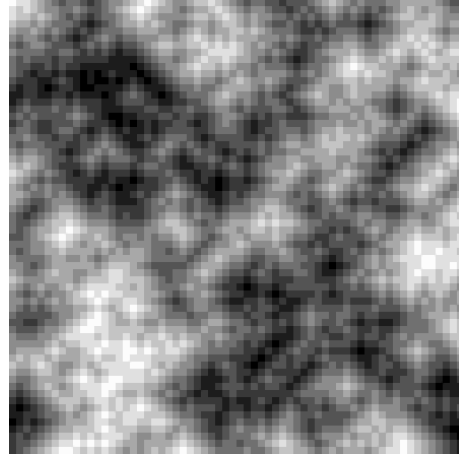


Figure 9. Collapsed in modulo addition order

My algorithm implemented a two-dimensional version of this modulo addition indexing which allows for much more organic results as can be seen in the figures above with minimal computational cost.

References

- M. Games. The fantastic combinations of john conway’s new solitaire game “life” by martin gardner. *Scientific American*, 223:120–123, 1970.
- B. Lindley. Iterating over the array in random order [duplicate]. <https://stackoverflow.com/a/18994414>, Sept 2013.
- S. Moreau. Generative design with the wave function collapse algorithm, Sep 2020. URL <https://www.bim42.com/2020/09/generative-design-with-the-wave-function-collapse-algorithm>.
- J. Orłowski and A. Lee. Parallel wave function collapse. <https://github.com/amylh/WaveCollapseGen>, 2019.
- N. Ru. Convolutions and the game of life, Dec 2017. URL <https://nicholasrui.com/2017/12/18/convolutions-and-the-game-of-life/>.